

Exercise #1

31. August 2022

Note. The exercises can be submitted in groups of three. All solutions have to be written into *one* Jupyter notebook (use Markdown cells for the mathematical answers). This notebook has to state the exercise number and *all group members* (again, up to 3) in the very beginning. One of the group members then submits this Jupyter notebook on Blackboard.

Problem 1.

- a) Prove that, for any nonsingular matrix $A \in \mathbb{R}^{n \times n}$,

$$\kappa_2(A) = \left(\frac{\lambda_{\max}}{\lambda_{\min}} \right)^{\frac{1}{2}},$$

where λ_{\min} is the smallest and λ_{\max} is the largest eigenvalue of the matrix $A^T A$.

- b) Show that the condition number $\kappa_2(Q)$ of an orthogonal matrix Q is equal to 1.
- c) Conversely, if $\kappa_2(A) = 1$ for the matrix A , show that all the eigenvalues of $A^T A$ are equal; deduce that A is a scalar multiple of an orthogonal matrix.

Problem 2.

We want to compute the LU factorisation with pivoting of a matrix A

$$PA = LU$$

where P is a permutation matrix, L is a lower-triangular matrix with unit diagonal, and U is an upper-triangular matrix. We represent the matrices in question as follows: The permutation matrix P is $n \times n$, but is represented as a vector P such that row number k in P is the canonical unit vector e_{P_k} . Let us illustrate this by an example

$$P = \begin{bmatrix} 3 \\ 1 \\ 2 \end{bmatrix} \Rightarrow P = \begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

We stipulate that a Python function takes a two-dimensional numpy-array A as input, and returns an *over-written* A which contains L and U in the following sense upon return:

$$\begin{aligned} A[P[i], j] &= L_{ij} \quad \text{for } i < j \\ A[P[i], j] &= U_{ij} \quad \text{for } i \geq j \end{aligned}$$

That L has 1 on the diagonal is always the case, so the diagonal of L needs not be stored. The remaining elements of L and U are zero and need not be stored either. The algorithm can be formulated as follows (compare to text book):

- Input: A of size $n \times n$.
- Initialisation:
 - Let $P_i = i$, $i = 0, \dots, n-1$ be a vector (array) with n components.
- for k in range (n-1):
 - a) Find index P_ℓ such that $|A_{P_\ell, k}| = \max_{k \leq i \leq n-1} |A_{P_i, k}|$, i.e. scan column k from the diagonal and down for the largest element in absolute value.
 - b) Swap P_k by P_ℓ .
 - c) Find multipliers $A_{P_i, k} \leftarrow \frac{A_{P_i, k}}{A_{P_k, k}}$, $i = k+1, \dots, n-1$.
 - d) Perform elimination, i.e. $A_{P_i, j} \leftarrow A_{P_i, j} - A_{P_i, k} \cdot A_{P_k, j}$, $i, j = k+1, \dots, n-1$

- Output: A,P.

There are – of course – implementations of these, often highly optimised for special cases (e.g. when A is sparse) but here we first want to learn how to code it ourselves. Let's also use this to our advantage.

- a) Write a function for LU-factorisation with row-wise pivoting as indicated above. A template could be

```
def mylu (A):
```

and it should return the pivot vector (permutation vector) P , and over-written version of A . You can also choose to copy A into some other matrix LU from the beginning using e.g.

```
LU = A.copy().
```

- b) Use the function `scipy.linalg.lu` to test your implementation from the first part. Write a test function `def mylutest(A)` that compares the result of your implementation to the one from SciPy. Call this function for example with

$$A = \begin{pmatrix} 2 & 5 & 8 & 7 \\ 5 & 2 & 2 & 8 \\ 7 & 5 & 6 & 6 \\ 5 & 4 & 4 & 8 \end{pmatrix}$$

(or `np.array ([[2, 5, 8, 7], [5, 2, 2, 8], [7, 5, 6, 6], [5, 4, 4, 8]])` to easier copy it).

- c) Test your function with a matrix A that does not meet our assumption of having full rank, for example by repeating a column. What happens?

Problem 3.

- a) Implement a

```
1 def forward_substitution(A,b):
```

function that takes a lower triangular matrix A and some vector b to solve $Ax = b$.

b) Implement a

```
1         def backward_substitution(A,b):
```

function that takes an upper triangular matrix A and some vector b to solve $Ax = b$ for this case.

c) Combine the last two parts with Problem 2 and implement a function

```
1         def my_solve(A,b):
```

for a square matrix A and a right hand side b that computes the LU decomposition of A and then uses the first two parts of this problem to compute the solution to $Ax = b$.

d) Let's compare our implementation for two cases as well as to the original. To be precise:

- fix some n , say $n=100$,
- generate a regular square matrix A (Hint: maybe create L and U here to be sure A is regular),
- generate m , say $m=200$ right hand sides b_k ,
- run an experiment where you time:
 - 1) calling m times `my_solve(A,b)` once for each b_k ,
 - 2) only computing the LU decomposition once and use backward and forward substitutions m times,
 - 3) using `np.linalg.solve`.

Where do the time differences from (1) to (2) come from and where the ones from (3) to (2)?